

# 03a\_CAMSHIFT

November 13, 2019

## 1 Assignment 3: CAMSHIFT

### 1.1 Paper

Read the paper “Bradski\_etal\_1998\_camshift.pdf” in KVV (under “Resources/papers”).

### 1.2 Calculate histogram

- Implement a function that creates a color histogram. Pass either an image and ROI, or the image underlying the ROI.
- For this purpose, a second (or third) parameter can be passed to specify the number of bins.
- Load the image “images/racecar.png” and convert the image to the HSV color space. Plot the Hue channel. **(RESULT)**

In [16]: *# dieser Code wurde als Musterlösung von Tobias Schülke zur Verfügung gestellt und von*

```
%matplotlib inline
from skimage import io, color
import numpy as np
from matplotlib import pyplot as plt
import matplotlib.patches as patches
import os
import warnings; warnings.simplefilter('ignore')
```

```
IMAGES_PER_ROW = 4
```

```
MIN_SATURATION_CAR = 0.2
MIN_VALUE_CAR = 0.5
MIN_SATURATION_TACO = 0.8
MIN_VALUE_TACO = 0.2
```

```
ROI_FRAME_MARGIN_CAR = 60
ROI_FRAME_MARGIN_TACO = 20
```

```
def create_color_histogram(img, bins=360, normalize=False):
    hues = color.rgb2hsv(img)[: , :, 0]
```

```

hist, bin_edges = np.histogram(hues, bins=bins, range=(0, 1))
if normalize:
    hist = hist / np.max(hist)
return hist, bin_edges

def plot_histogram(ax, hist, bins=360):
    X = range(0, bins)
    hsv_colors = np.array([[h/bins, 1, 1] for h in X])
    rgb_colors = (
        color.hsv2rgb(
            # reshape to make 'hsv2rgb' work because it expects an image
            hsv_colors.reshape((bins, 1, 3))
        )
        .reshape((bins, 3))
    )
    return ax.bar(X, hist, color=rgb_colors)

def create_and_plot_histogram(ax, img, bins=360):
    hist, bin_edges = create_color_histogram(img, bins=bins)
    return hist, bin_edges, plot_histogram(ax, hist, bins=bins)


image = io.imread('images/racecar.png')

fig = plt.figure(figsize=(18, 12))
ax1 = plt.subplot(2, 2, 1)
ax2 = plt.subplot(2, 2, 2)

ax1.imshow(image)
ax2.imshow(color.rgb2hsv(image)[: , : , 0])

```

None



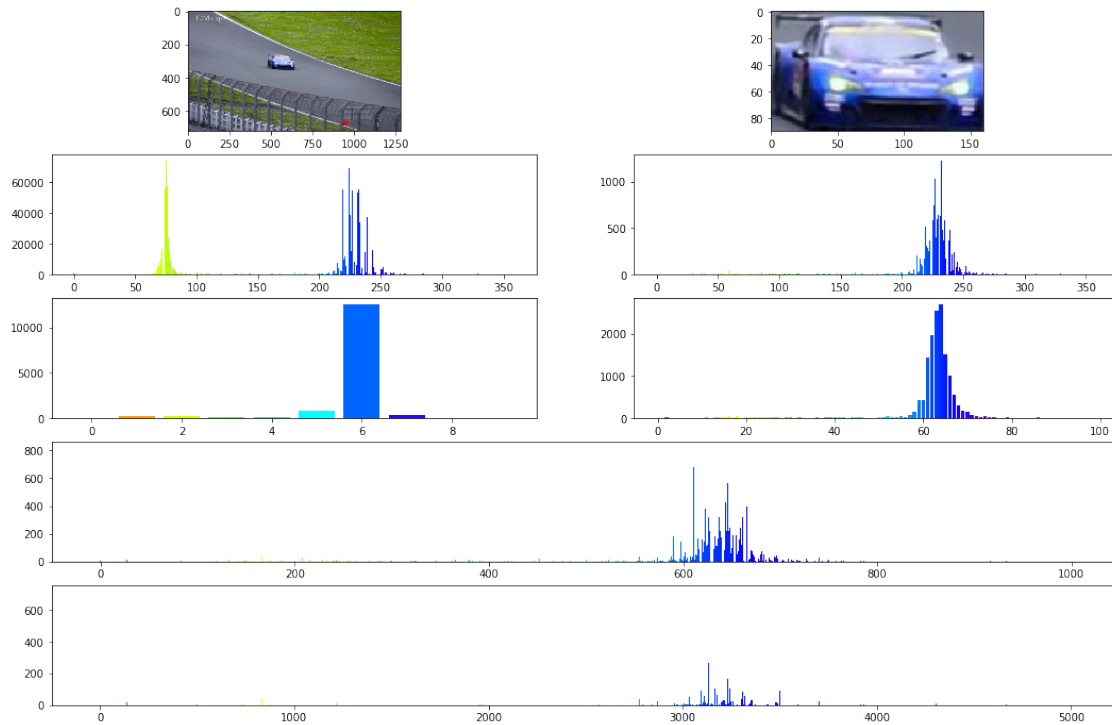
- display the histogram for the Hue channel for the entire image and for the RIO  $(x,y) = (480, 260)$  to  $(640, 350)$ . Vary the number of bins for testing purposes (**RESULT**).

```
In [17]: image_car = image[260:350, 480:640]
```

```
fig = plt.figure(figsize=(18, 12))
grid = fig.add_gridspec(5, 2)
ax1 = fig.add_subplot(grid[0, 0])
ax2 = fig.add_subplot(grid[0, 1])
ax3 = fig.add_subplot(grid[1, 0])
ax4 = fig.add_subplot(grid[1, 1])
ax5 = fig.add_subplot(grid[2, 0])
ax6 = fig.add_subplot(grid[2, 1])
ax7 = fig.add_subplot(grid[3, :])
ax8 = fig.add_subplot(grid[4, :])

ax1.imshow(image)
ax2.imshow(image_car)
create_and_plot_histogram(ax3, image)
create_and_plot_histogram(ax4, image_car)
create_and_plot_histogram(ax5, image_car, bins=10)
create_and_plot_histogram(ax6, image_car, bins=100)
create_and_plot_histogram(ax7, image_car, bins=1000)
create_and_plot_histogram(ax8, image_car, bins=5000)
```

None




### 1.3 probability distribution

- implement the method outlined in the CAMSHIFT paper to create a probability distribution for a given object's hue histogram
- create the histogram of the car from the image "racecar.png" and apply the new function to the last frame of the video (images/racecar/151.jpg) (**RESULT**)

```
In [117]: # Tip: in der Nacht sind alle Katzen grau ;)
def create_prob_distribution(img, hist, bin_edges, min_saturation, min_value):
    img_hsv = color.rgb2hsv(img)
    # We want grayscale.
    prob_dist = np.zeros((img.shape[0], img.shape[1]))
    # print('hsv', img_hsv.shape)
    # Ignore irrelevant pixels (see page 3 in the paper).
    indices = np.logical_and(img_hsv[:, :, 1] >= min_saturation, img_hsv[:, :, 2] >=
    # print('indices', indices.shape)
    hues = img_hsv[indices, 0]
    # print('hues', hues.shape)
    binned_hues = np.digitize(hues, bin_edges)
    # print('>', binned_hues.shape)
    max_index = len(bin_edges) - 1
    binned_hues[np.where(binned_hues == max_index)] -= 1
    prob_dist[indices] = hist[binned_hues]
    return prob_dist

last_frame = io.imread('images/racecar/151.jpeg')
histogram_car, bin_edges = create_color_histogram(image_car, bins=100, normalize=True)
prob_dist = create_prob_distribution(last_frame, histogram_car, bin_edges, MIN_SATUR
io.imshow(prob_dist)

None
```




### 1.3.1 show ROI in image

```
In [19]: def draw_roi(image, x, y, width, height, out=plt):
          ax = plt.gca()
          ax.add_patch(
              patches.Rectangle(
                  (x, y),
                  width,
                  height,
                  fill=False,
                  edgecolor="red"
              )
          )

          out.imshow(image)

draw_roi(image, 480, 260, 160, 90)
```



## 1.4 Exercise 3.2 - Mean Shift

- Implement the Mean Shift method for a ROI as described in the lecture. Test the algorithm on the image sequences *"images/racecar/.jpg"* or *"images/taco/.jpg"*. Select the appropriate tracking window (to reduce the file size I have reduced the images by a factor of 2, i.e. the ROI from above must be adjusted accordingly).
- Draw the trajectory of the objects as returned by Mean Shift. (**RESULT**)

In [217]: `from math import ceil`

```
from numpy.linalg import norm
from skimage.draw import circle
```

```
def get_center(rect):
    return np.mean(rect.T, axis=1).round()
```

```
def get_diag_len(rect):
    top_left, bottom_right = rect
    return np.sqrt(np.square(top_left[0] - bottom_right[0]) + np.square(top_left[1] - bottom_right[1]))
```

```
def get_roi(img, rect):
    top_left, bottom_right = rect
    return img[top_left[0]:bottom_right[0], top_left[1]:bottom_right[1]]
```

```

def get_shifted_next_window(window, moments, v, roi):
    return window + v.round().astype(np.int)

def mean_shift(initial_window, frames, create_prop_dist,
               *,
               frame_indices_to_draw=None,
               keyframes=None,
               ellipsis_indices_to_draw = None,
               shift_threshold=4,
               max_iterations=5,
               get_moments=get_moments,
               get_next_window=get_shifted_next_window,
               draw_ellipsis=None,):
    first_frame, frames = frames[0], frames[1:]

    if frame_indices_to_draw is None:
        frame_indices_to_draw = []
    if keyframes is None:
        keyframes = []
    if ellipsis_indices_to_draw is None:
        ellipsis_indices_to_draw = []

    window = initial_window
    centroids = [get_center(window)]

    NCOLS = 3
    fig, axes = plt.subplots(
        nrows=ceil((len(frame_indices_to_draw) + len(keyframes)) / NCOLS),
        ncols=NCOLS,
    )
    fig.set_figheight(18)
    fig.set_figwidth(18)
    plt.subplots_adjust(wspace = .01)
    fig.tight_layout()
    axes = np.asarray(axes).ravel()

    i = -1
    plot_index = 0
    for frame in frames:
        i += 1
        prev_centroid = centroids[-1]
        temp_centroids = [prev_centroid]

        num_iterations = -1
        while True:
            num_iterations += 1

```

```

roi = get_roi(frame, window)
hist, bin_edges = create_color_histogram(roi, bins=100, normalize=True)
prop_dist = create_prop_dist(roi, hist, bin_edges)

moments = get_moments(prop_dist)
if moments[0] == 0:
    print('M_00 == 0, i =', i)
    break

centroid = window[0] + get_centroid(moments)
temp_centroids.append(centroid)

v = centroid - prev_centroid
prev_centroid = centroid
# window += v.round().astype(np.int)
window = get_next_window(window, moments, v, roi)
if norm(v) <= shift_threshold or num_iterations > max_iterations:
    break

centroids.append(centroid)
if i in frame_indices_to_draw:
    axes[plot_index].imshow(draw_centroids(frame, temp_centroids))
    axes[plot_index].set_title(f'shifting mean {i}')
    plot_index += 1
if i in keyframes:
    j = keyframes.index(i)
    start = keyframes[j - 1] if j > 0 else 0
    c = centroids[start:keyframes[j]]
    axes[plot_index].imshow(draw_centroids(frame, c))
    axes[plot_index].set_title(f'frames {start} - {keyframes[j]}')
    plot_index += 1
if i in ellipsis_indices_to_draw and draw_ellipsis:
    axes[plot_index].imshow(draw_ellipsis(frame, centroid, moments))
    axes[plot_index].set_title(f'roll {i}')
    plot_index += 1

return centroids

def get_moments(prop_dist):
    """Returns 0th and 1st image moments."""
    M_00 = np.sum(prop_dist)
    # Because we only have positive values (propabilities)
    # we know that all values must be zero, thus:
    #  $M_{00} == x \ y \ I(x,y) == 0 \Rightarrow x \ y \ x \ I(x,y) == M_{10} == 0$ 
    if M_00 > 0:
        num_rows, num_cols = prop_dist.shape
        x_indices = np.arange(0, num_cols)

```



```

        y_indices = np.arange(0, num_rows)
        M_10 = np.sum(prop_dist[:] * x_indices)
        M_01 = np.sum(prop_dist.T[:] * y_indices)
        M_11 = np.sum((prop_dist[:] * x_indices).T[:] * y_indices)
    else:
        M_10 = M_01 = M_11 = 0
    return (M_00, (M_10, M_01, M_11))

def get_centroid(moments, row_first=True):
    M_00, (M_10, M_01, M_11) = moments[0], moments[1]
    centroid = np.array([M_01, M_10])
    if M_00 > 0:
        centroid = centroid / M_00
    if row_first:
        return centroid
    return np.flip(centroid)

def draw_centroids(img, centroids):
    img = np.copy(img)
    for centroid in centroids:
        rr, cc = circle(*centroid, radius=5, shape=img.shape)
        img[rr, cc] = [255, 0, 0]
    return img

frames = io.imread_collection('images/racecar/*')
initial_window = np.array([[130, 237], [175, 316]])

centroids = mean_shift(
    initial_window,
    frames,
    lambda img, hist, bin_edges:
        create_prob_distribution(img, hist, bin_edges, MIN_SATURATION_CAR, MIN_VALUE,
        frame_indices_to_draw=[2, 30],
        keyframes=[20, 60, 64, 71, 89, 111, 150],
)

```

None

```

M_00 == 0, i = 89
M_00 == 0, i = 90
M_00 == 0, i = 92

```



## 1.5 Exercise 3.3 - CAMSHIFT

- extend your algorithm by adjusting the size of the ROI and finding the object's orientation
- execute the algorithm again on one of the image sequences and draw an ellipse on the image, which represents the found parameters (**RESULT**)

```
In [203]: def get_scaled_next_window(window, moments, v, roi):
          """
          For 2D color probability distributions where the maximum pixel value is 255,
          we set window size  $s$  to
           $s = 2 * \sqrt{M_{00} / 256}$ 
          """
          shifted_window = get_shifted_next_window(window, moments, v, roi)
          M_00 = moments[0]
          max_hue = np.max(roi[:, :, 0])
          scaling_factor = 2 * np.sqrt(M_00 / max_hue)
          scaled_window = shifted_window
```

```

        return scaled_window


frames = io.imread_collection('images/racecar/*')
initial_window = np.array([[130, 237], [175, 316]])

centroids = mean_shift(
    initial_window,
    frames,
    lambda img, hist, bin_edges:
        create_prob_distribution(img, hist, bin_edges, MIN_SATURATION_CAR, MIN_VALUE,
        frame_indices_to_draw=[*range(0, 9)],
        keyframes=[],
        get_next_window=get_scaled_next_window,
)

None

M_00 == 0, i = 89
M_00 == 0, i = 90
M_00 == 0, i = 92

```



```
In [221]: from skimage.draw import ellipse_perimeter
```

```
def get_moments2(prop_dist):
    """Returns 0th, 1st and 2nd order image moments."""
    # (M_00, (M_10, M_01))
    lower_order_moments = M_00, M_1 = get_moments(prop_dist)

    if M_00 > 0:
        num_rows, num_cols = prop_dist.shape
        x_indices = np.arange(0, num_cols)
        y_indices = np.arange(0, num_rows)
        M_20 = np.sum(prop_dist[:] * np.square(x_indices))
        M_02 = np.sum(prop_dist.T[:] * np.square(y_indices))
    else:
        M_20 = M_02 = 0
    return (*lower_order_moments, (M_20, M_02))
```

```

def get_roll_params(centroid, moments):
    x_c, y_c = centroid
    M_00, (M_10, M_01, M_11), (M_20, M_02) = moments

    if M_00 > 0:
        a = M_20/M_00 - x_c**2
        b = 2 * (M_11/M_00 - x_c*y_c)
        c = M_02/M_00 - y_c**2
    else:
        a = -x_c**2
        b = 2 * (-x_c*y_c)
        c = -y_c**2
    p = a + c
    q = np.sqrt(b**2 + (a - c)**2)

    angle = 0.5 * np.arctan(b / (a - c))

    l = np.sqrt(0.5 * (p + q))
    w = np.sqrt(0.5 * (p - q))
    return angle, l, w

def draw_ellipsis(img, centroid, moments):
    img = np.copy(img)
    orientation, l, w = get_roll_params(centroid, moments)
    rr, cc = ellipse_perimeter(*centroid, int(l), int(w), orientation, shape=img.shape)
    img[rr, cc] = [255, 0, 0]
    return img

frames = io.imread_collection('images/racecar/*')
initial_window = np.array([[130, 237], [175, 316]])

centroids = mean_shift(
    initial_window,
    frames,
    lambda img, hist, bin_edges:
        create_prob_distribution(img, hist, bin_edges, MIN_SATURATION_CAR, MIN_VALUE),
    frame_indices_to_draw=[*range(0, 9)],
    keyframes=[],
    ellipsis_indices_to_draw=[*range(0, 9)],
    get_moments=get_moments2,
    get_next_window=get_scaled_next_window,
    draw_ellipsis=draw_ellipsis,
)

```

None

-----  
ValueError

Traceback (most recent call last)

```
<ipython-input-221-98baf73d479f> in <module>()
    62     get_moments=get_moments2,
    63     get_next_window=get_scaled_next_window,
--> 64     draw_ellipsis=draw_ellipsis,
    65 )
    66

<ipython-input-217-8f8fcd6819e4> in mean_shift(initial_window, frames, create_prop_dis
    99         plot_index += 1
   100         if i in ellipsis_indices_to_draw and draw_ellipsis:
--> 101             axes[plot_index].imshow(draw_ellipsis(frame, centroid, moments))
   102             axes[plot_index].set_title(f'roll {i}')
   103             plot_index += 1

<ipython-input-221-98baf73d479f> in draw_ellipsis(img, centroid, moments)
    43     img = np.copy(img)
    44     orientation, l, w = get_roll_params(centroid, moments)
--> 45     rr, cc = ellipse_perimeter(*centroid, int(l), int(w), orientation, shape=img.sh
    46     img[rr, cc] = [255, 0, 0]
    47     return img
```

ValueError: cannot convert float NaN to integer

